

Introduction  
oooooooo

MDP & solution  
oooooooooooooo

Dynamic programming  
oooooooooooooo

Monte Carlo  
ooooooo

Exploration  
oooo

Temporal differences  
oooooooooooooo

Improvements  
oooooooo

# Basics of reinforcement learning

Lucian Buşoniu

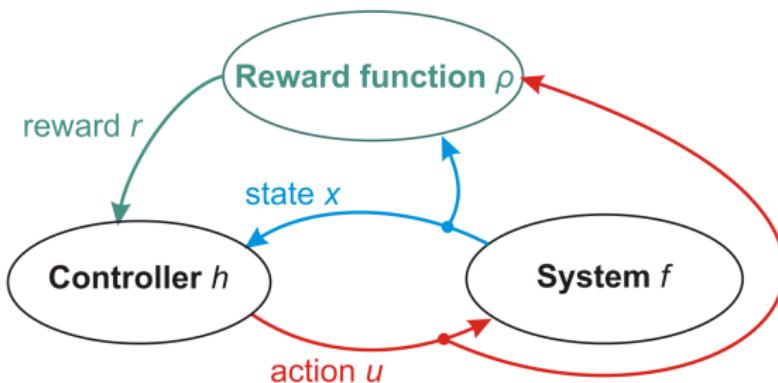
TMLSS, 20 July 2018



# Main idea of reinforcement learning (RL)

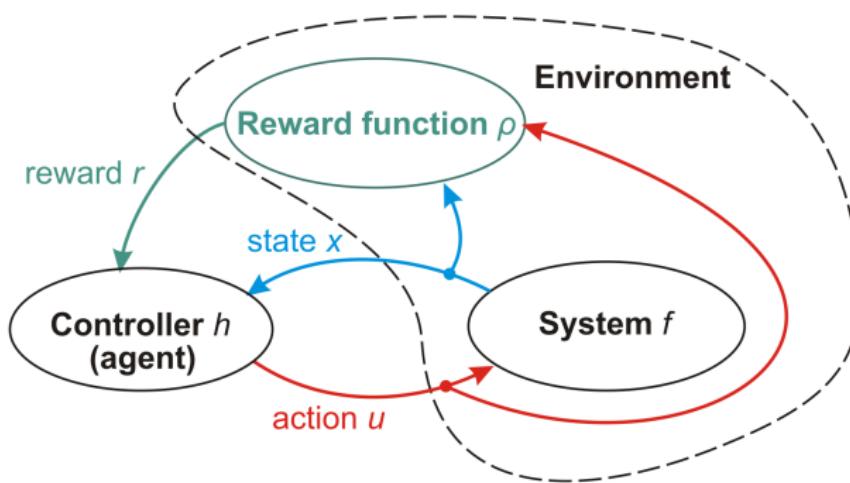
Learn a sequential decision policy  
to optimize the cumulative performance  
of an unknown system

## RL principle



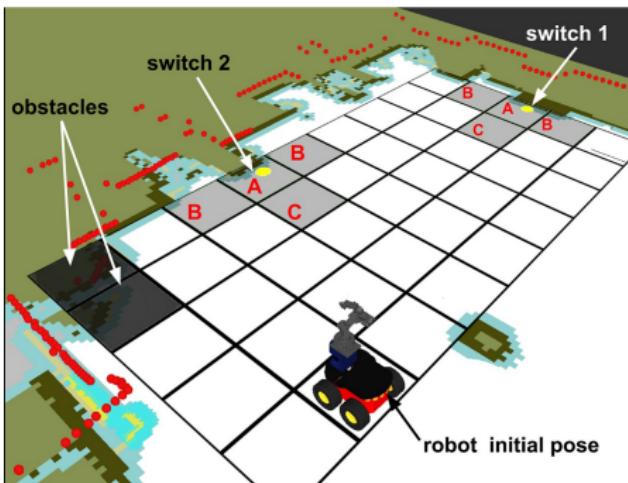
- Interact with system: measure **states**, apply **actions**
  - Performance feedback in the form of **rewards**
  - Inspired by human and animal learning

# RL principle: AI view



- Agent embedded in an environment that feeds back states and rewards

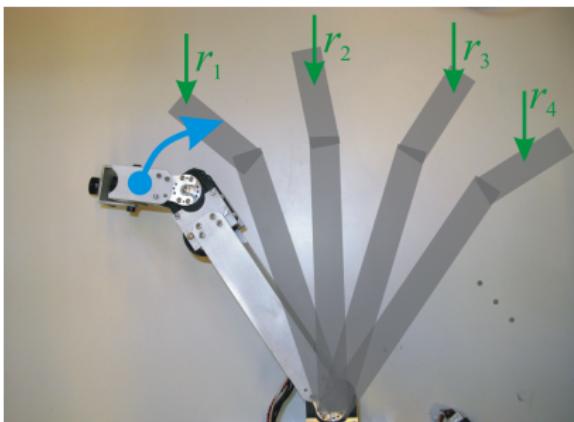
## Example: Domestic robot



A domestic robot ensures light switches are off  
Abstractization to high-level control (physical actions implemented by low-level controllers)

- **States**: grid coordinates, switch states
  - **Actions**: movements NSEW, toggling switch
  - **Rewards**: when switches toggled on→off

# Example: Robot arm

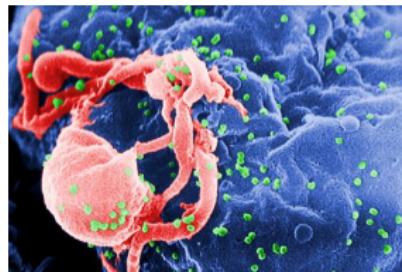


## Low-level control

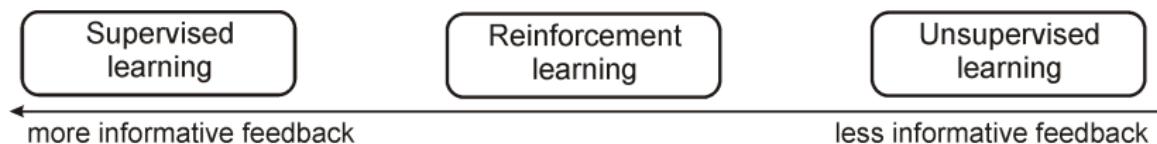
- **States:** link angles and angular velocities
- **Actions:** motor voltages
- **Rewards:** e.g. to reach a desired configuration, give larger rewards as robot gets closer to it

# Many other applications

Artificial intelligence, control, medicine, multiagent systems, economics etc.



# RL on the machine learning spectrum

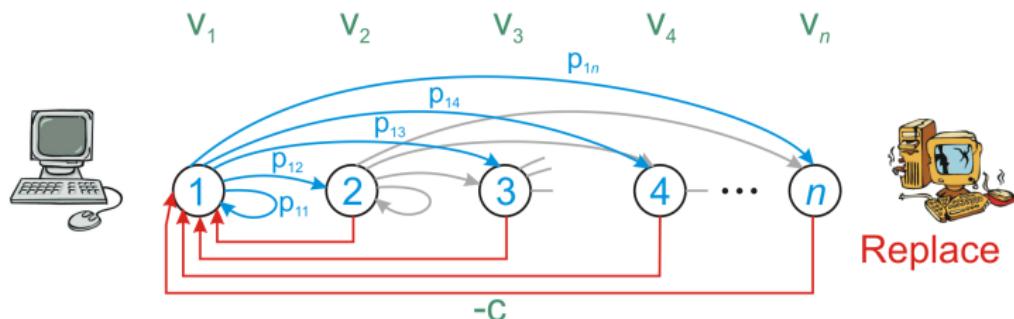


- Supervised: for each training sample, **correct output** known
- Unsupervised: only input samples, **no outputs**; find patterns in the data
- Reinforcement: correct actions not available, **only rewards**

But note: RL finds **dynamical optimal control!**

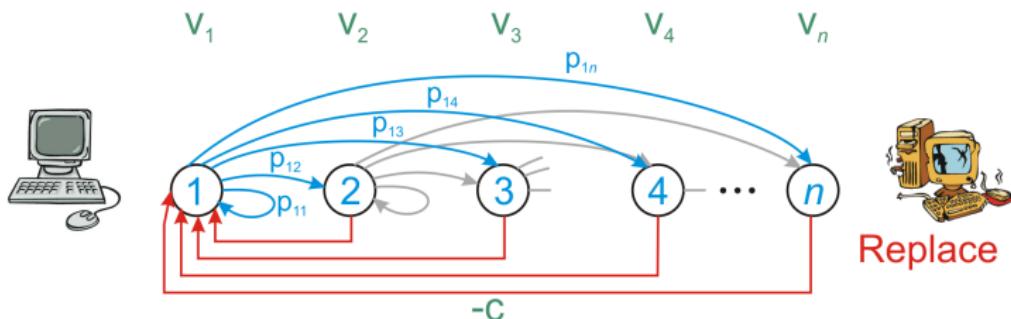
- 1 Introduction
- 2 Markov decision process & optimal solution
- 3 Dynamic programming, DP
- 4 Monte Carlo, MC
- 5 Exploration basics
- 6 Temporal differences, TD
- 7 Improving data efficiency

## Example: Machine replacement



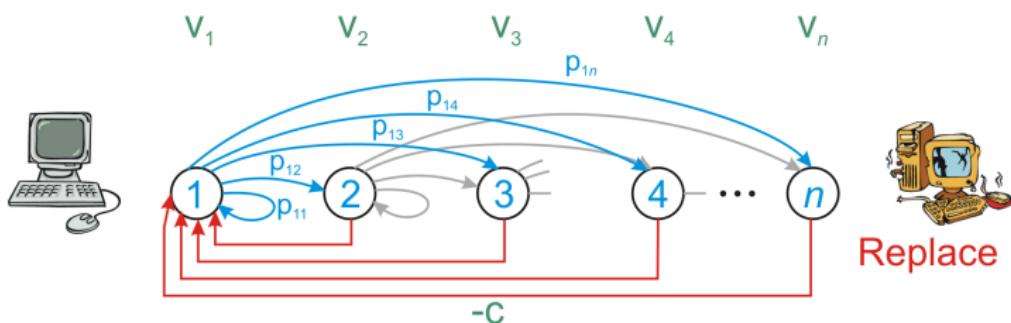
- Consider machine with  $n$  different wear levels  
 $1 =$  perfect working order,  $n =$  fully degraded
- Produces revenue  $v_i$  when operating at level  $i$
- Stochastic wear: level  $i$  increases to  $j > i$  with probas  $p_{ij}$ , stays  $i$  with  $p_{ii} = 1 - p_{i,i+1} - \dots - p_{i,n}$
- Machine can be replaced at any time (assumed instantaneously), paying cost  $c$

# Machine replacement: States and actions



- State  $x = \text{wear level}$ ,  
state space  $X = \{1, 2, \dots, n\}$
- Action  $u = \text{whether to Wait or Replace}$ ,  
action space:  $U = \{W, R\}$

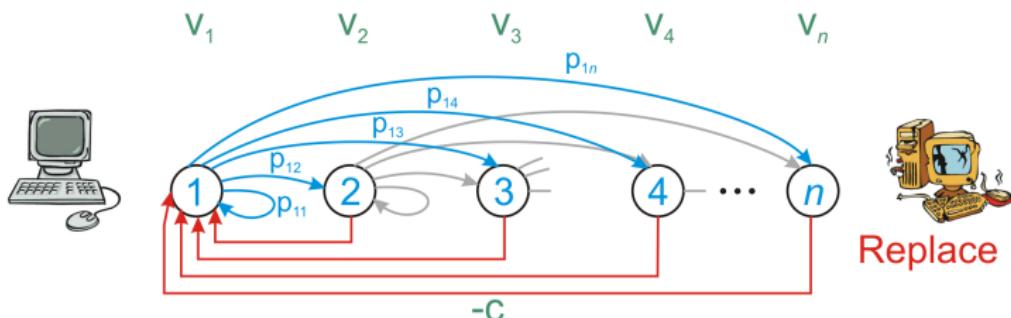
# Machine replacement: Transition function



Transition function  $f(x, u, x')$  gives the probability of reaching state  $x'$  after applying action  $u$  in state  $x$ :

$$f(x = i, u, x' = j) = \begin{cases} p_{ij} & \text{if } u = W \text{ and } i \leq j \\ 1 & \text{if } u = R \text{ and } j = 1 \\ 0 & \text{otherwise} \end{cases}$$

# Machine replacement: Reward function



Reward function  $\rho(x, u, x')$  gives reward resulting from transitioning from  $x$  to  $x'$  after applying  $u$ :

$$\rho(x = i, u, x' = j) = \begin{cases} v_i & \text{if } u = W \\ -c + v_1 & \text{if } u = R \end{cases}$$

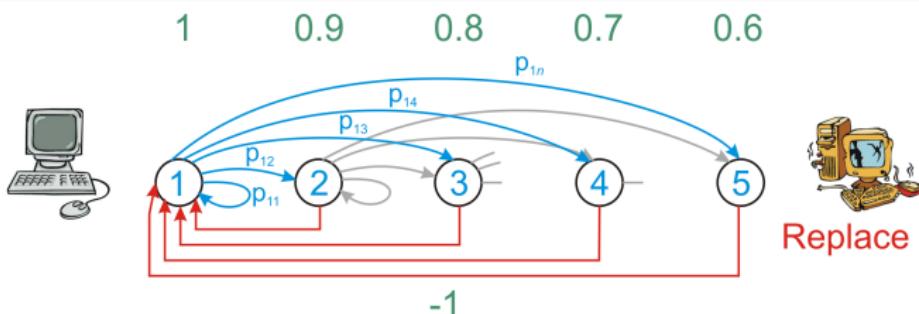
# General case: Markov decision process

## Markov decision process (MDP)

- ① State space  $X$
- ② Action space  $U$
- ③ Transition function  $f(x, u, x')$ ,  $f : X \times U \times X \rightarrow [0, 1]$
- ④ Reward function  $\rho(x, u, x')$ ,  $\rho : X \times U \times X \rightarrow \mathbb{R}$

Some MDPs have terminal states (e.g. success, failure), that once reached cannot be left and provide no additional reward

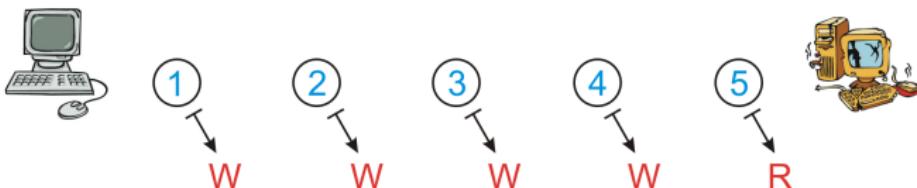
# Machine replacement: Specific example



- $n = 5$  wear levels
- Revenue:  $v_1 = 1, v_2 = 0.9, \dots, v_5 = 0.6$
- Cost of new machine:  $c = 1$
- Wear increase probabilities:

$$[p_{ij}] = \begin{bmatrix} 0.6 & 0.3 & 0.1 & 0 & 0 \\ 0 & 0.6 & 0.3 & 0.1 & 0 \\ 0 & 0 & 0.6 & 0.3 & 0.1 \\ 0 & 0 & 0 & 0.7 & 0.3 \\ 0 & 0 & 0 & 0 & 1.0 \end{bmatrix}$$

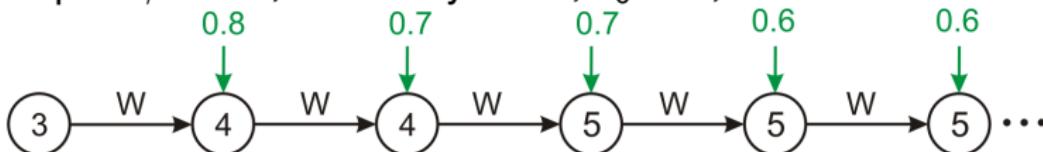
# Control policy



- **Control policy**  $h : X \rightarrow U$ : maps states  $x$  to actions  $u$
- Example for machine replacement:  $h(1) = \dots = h(4) = W$ ,  $h(5) = R$

# Return and objective

Example:  $\gamma = 0.9$ ,  $h$  = always wait,  $x_0 = 4$ , and trial:



$$\begin{aligned}\gamma^0 r_1 + \gamma^1 r_2 + \gamma^2 r_3 + \gamma^3 r_4 + \dots &= 0.9^0 \cdot 0.8 + 0.9^1 \cdot 0.7 + \\ &+ 0.9^2 \cdot 0.7 + 0.9^3 \cdot 0.6 + 0.9^4 \cdot 0.6 + \dots = 6.3710\end{aligned}$$

## Objective

Find  $h$  that maximizes from any  $x_0$  the **expected return** under the stochastic transitions:

$$R^h(x_0) = E_{x_1, x_2, \dots} \left\{ \sum_{k=0}^{\infty} \gamma^k \rho(x_k, h(x_k), x_{k+1}) \right\}$$

Note: There are also other types of return!

# Discount factor

Discount factor  $\gamma \in [0, 1)$ :

- represents an increasing uncertainty about the future
- bounds the infinite sum (assuming rewards bounded)
- helps the convergence of algorithms

To choose  $\gamma$ , **trade-off** between:

- ➊ Long-term quality of the solution (large  $\gamma$ )
- ➋ Simplicity of the problem (small  $\gamma$ )

In practice,  $\gamma$  should be sufficiently large so as not to ignore important later rewards



# Q-function

**Q-function** of a policy  $h$  is the expected return achieved by executing  $u_0$  in  $x_0$  and then following  $h$

$$Q^h(x_0, u_0) = \mathbb{E}_{x_1} \left\{ \rho(x_0, u_0, x_1) + \gamma R^h(x_1) \right\}$$

$Q^h$  measures the quality of state-action pairs

# Bellman equation of $Q^h$

Go one step further in the equation:

$$\begin{aligned} Q^h(x_0, u_0) &= \mathbb{E}_{x_1} \left\{ \rho(x_0, u_0, x_1) + \gamma R^h(x_1) \right\} \\ &= \mathbb{E}_{x_1} \left\{ \rho(x_0, u_0, x_1) + \gamma \mathbb{E}_{x_2} \left\{ \rho(x_1, h(x_1), x_2) + \gamma R^h(x_2) \right\} \right\} \\ &= \mathbb{E}_{x_1} \left\{ \rho(x_0, u_0, x_1) + \gamma Q^h(x_1, h(x_1)) \right\} \end{aligned}$$

⇒ **Bellman equation for  $Q^h$**

$$\begin{aligned} Q^h(x, u) &= \mathbb{E}_{x'} \left\{ \rho(x, u, x') + \gamma Q^h(x', h(x')) \right\} \\ &= \sum_{x'} f(x, u, x') \left[ \rho(x, u, x') + \gamma Q^h(x', h(x')) \right] \end{aligned}$$

# Optimal solution and Bellman optimality equation

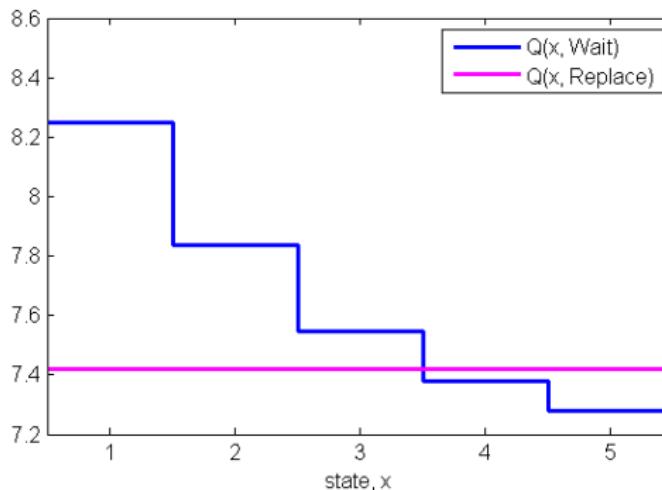
- **Optimal Q-function:**  $Q^* = \max_h Q^h$
- ⇒ “Greedy” policy in  $Q^*$ :  $h^*(x) = \arg \max_u Q^*(x, u)$   
is **optimal**, i.e. achieves maximal returns  
(if multiple actions maximize, break ties arbitrarily)

## Bellman optimality equation (for $Q^*$ )

$$\begin{aligned} Q^*(x, u) &= E_{x'} \left\{ \rho(x, u, x') + \gamma \max_{u'} Q^*(x', u') \right\} \\ &= \sum_{x'} f(x, u, x') \left[ \rho(x, u, x') + \gamma \max_{u'} Q^*(x', u') \right] \end{aligned}$$

# Machine replacement: Optimal solution

Discount factor  $\gamma = 0.9$



Introduction  
oooooooo

MDP & solution  
oooooooooooooo

Dynamic programming  
oooooooooooooo

Monte Carlo  
ooooooo

Exploration  
oooo

Temporal differences  
oooooooooooooo

Improvements  
oooooooo

Up next:

Algorithms to find the optimal solution



# Algorithm landscape

By model usage:

- **Model-based**:  $f, \rho$  known a priori
- **Model-free**:  $f, \rho$  unknown
- **Model-learning**:  $f, \rho$  learned from data

Model-based usually called dynamic programming (DP);  
needed as a stepping stone to RL

By interaction level:

- **Offline**: algorithm runs in advance
- **Online**: algorithm runs with the system

We focus on **exact** case:  $x, u$  small number of discrete values,  
so we can exactly represent solutions. In practice, function  
approximation often needed – covered in Doina's talk



- 1 Introduction
- 2 Markov decision process & optimal solution
- 3 Dynamic programming, DP
- 4 Monte Carlo, MC
- 5 Exploration basics
- 6 Temporal differences, TD
- 7 Improving data efficiency

# Q-iteration

Transforms Bellman optimality equation:

$$Q^*(x, u) = \sum_{x'} f(x, u, x') \left[ \rho(x, u, x') + \gamma \max_{u'} Q^*(x', u') \right]$$

into an **iterative procedure**:

## Q-iteration

**repeat** at each iteration  $\ell$

**for all**  $x, u$  **do**

$$Q_{\ell+1}(x, u) \leftarrow \sum_{x'} f(x, u, x') [\rho(x, u, x') + \gamma \max_{u'} Q_\ell(x', u')]$$

**end for**

**until** convergence to  $Q^*$

Once  $Q^*$  available:  $h^*(x) = \arg \max_u Q^*(x, u)$

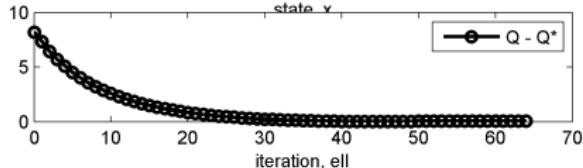
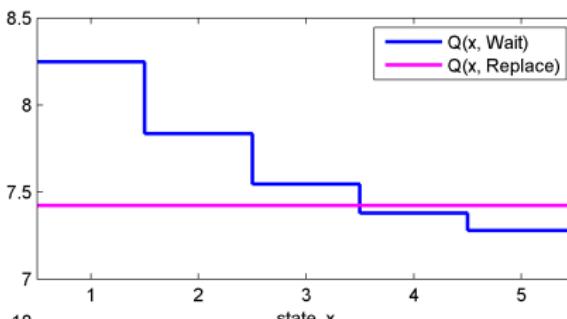
Q-iteration belongs to the class of value iteration algorithms



# Machine replacement: Q-iteration demo

Discount factor  $\gamma = 0.9$

Q-iteration, ell=64



# Machine replacement: Q-iteration demo

$$Q_{\ell+1}(x, u) \leftarrow \sum_{x'} f(x, u, x') [\rho(x, u, x') + \gamma \max_{u'} Q_\ell(x', u')]$$

	$x = 1$	$x = 2$	$x = 3$	$x = 4$	$x = 5$
$Q_0$	0 ; 0	0 ; 0	0 ; 0	0 ; 0	0 ; 0
$Q_1$	1 ; 0	0.9 ; 0	0.8 ; 0	0.7 ; 0	0.6 ; 0
$Q_2$	1.86 ; 0.9	1.67 ; 0.9	1.48 ; 0.9	1.3 ; 0.9	1.14 ; 0.9
$Q_3$	2.58 ; 1.67	2.31 ; 1.67	2.05 ; 1.67	1.83 ; 1.67	1.63 ; 1.67
$Q_4$	3.2 ; 2.33	2.87 ; 2.33	2.55 ; 2.33	2.3 ; 2.33	2.1 ; 2.33
...	...	...	...	...	...
$Q_{64}$	8.25 ; 7.42	7.84 ; 7.42	7.55 ; 7.42	7.38 ; 7.42	7.28 ; 7.42
$Q_{65}$	8.25 ; 7.42	7.84 ; 7.42	7.55 ; 7.42	7.38 ; 7.42	7.28 ; 7.42
$h^*$	W	W	W	R	R

$$h^*(x) = \arg \max_u Q^*(x, u)$$

# Policy iteration

## Policy iteration

initialize policy  $h_0$

**repeat** at each iteration  $\ell$

1: **policy evaluation:** find  $Q^{h_\ell}$

2: **policy improvement:**

$$h_{\ell+1}(x) \leftarrow \arg \max_u Q^{h_\ell}(x, u)$$

**until** convergence to  $h^*$

# Iterative policy evaluation

Similarly to Q-iteration, transforms Bellman equation for  $Q^h$ :

$$Q^h(x, u) = \sum_{x'} f(x, u, x') \left[ \rho(x, u, x') + \gamma Q^h(x', h(x')) \right]$$

into an iterative procedure:

## Iterative policy evaluation

**repeat** at each iteration  $\tau$

**for all**  $x, u$  **do**

$$Q_{\tau+1}(x, u) \leftarrow \sum_{x'} f(x, u, x') \left[ \rho(x, u, x') + \gamma Q_\tau(x', h(x')) \right]$$

**end for**

**until** convergence to  $Q^h$

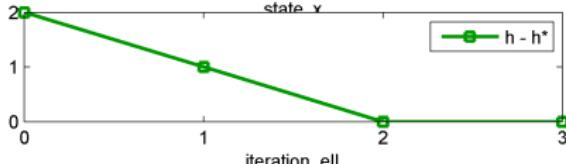
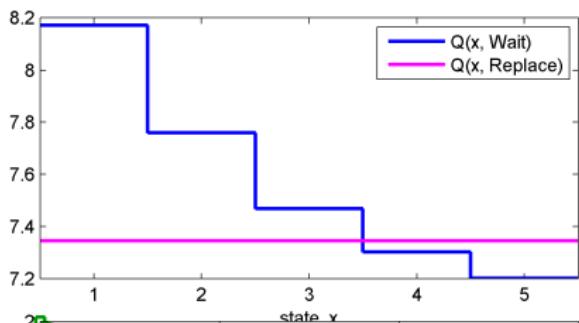
(other options exist, e.g. solving the linear system)



# Machine replacement: policy iteration demo

Discount factor  $\gamma = 0.9$

Policy iteration, ell=3



# Machine replacement: policy iteration

$$Q_{\tau+1}(x, u) \leftarrow \sum_{x'} f(x, u, x') [\rho(x, u, x') + \gamma Q_\tau(x', h(x'))]$$

$$h_{\ell+1}(x) \leftarrow \arg \max_u Q^{h_\ell}(x, u)$$

	$x = 1$	$x = 2$	$x = 3$	$x = 4$	$x = 5$
$h_0$	W	W	W	W	W
$Q_0$	0 ; 0	0 ; 0	0 ; 0	0 ; 0	0 ; 0
$Q_1$	1 ; 0	0.9 ; 0	0.8 ; 0	0.7 ; 0	0.6 ; 0
$Q_2$	1.86 ; 0.9	1.67 ; 0.9	1.48 ; 0.9	1.3 ; 0.9	1.14 ; 0.9
$Q_3$	2.58 ; 1.67	2.31 ; 1.67	2.05 ; 1.67	1.83 ; 1.67	1.63 ; 1.67
...	...	...	...	...	...
$Q_{39}$	7.51 ; 6.75	6.95 ; 6.75	6.49 ; 6.75	6.17 ; 6.75	5.9 ; 6.75
$Q_{40}$	7.52 ; 6.75	6.96 ; 6.75	6.5 ; 6.75	6.18 ; 6.75	5.91 ; 6.75
$h_1$	W	W	R	R	R

...algorithm continues...



# Machine replacement: policy iteration (cont'd)

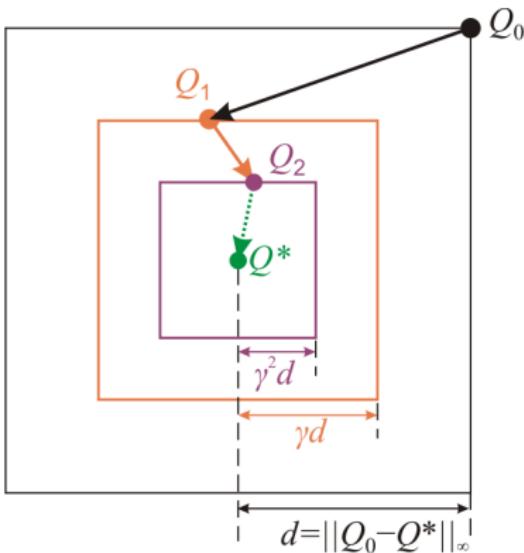
	$x = 1$	$x = 2$	$x = 3$	$x = 4$	$x = 5$
$h_1$	W	W	R	R	R
$Q_0$	0 ; 0	0 ; 0	0 ; 0	0 ; 0	0 ; 0
...	...	...	...	...	...
$Q_{43}$	8.01 ; 7.2	7.57 ; 7.2	7.27 ; 7.2	7.17 ; 7.2	7.07 ; 7.2
$h_2$	W	W	W	R	R
$Q_0$	0 ; 0	0 ; 0	0 ; 0	0 ; 0	0 ; 0
...	...	...	...	...	...
$Q_{43}$	8.17 ; 7.35	7.76 ; 7.35	7.47 ; 7.35	7.3 ; 7.35	7.2 ; 7.35
$h_3$	W	W	W	R	R

# Convergence of Q-iteration

- Each iteration a contraction with factor  $\gamma$  in  $\infty$ -norm:

$$\|Q_{\ell+1} - Q^*\|_\infty \leq \gamma \|Q_\ell - Q^*\|_\infty$$

⇒ Q-iteration **monotonically converges** to  $Q^*$ ,  
**with convergence rate**  $\gamma \Rightarrow \gamma$  helps convergence



# Stopping condition of Q-iteration

- Convergence to  $Q^*$  only guaranteed asymptotically, as  $\ell \rightarrow \infty$
- In practice, algorithm can be stopped when:

$$\|Q_{\ell+1} - Q_\ell\| \leq \varepsilon_{\text{qiter}}$$

# Convergence of policy iteration

Policy evaluation component – like Q-iteration:

- Iterative policy evaluation contraction with factor  $\gamma$
- ⇒ **monotonic convergence** to  $Q^h$ , with rate  $\gamma$

Complete policy iteration algorithm:

- Policy is either improved or already optimal
  - But the maximum number of improvements is finite! ( $|U|^{|X|}$ )
- ⇒ **convergence** to  $h^*$  in a finite number of iterations

# Stopping conditions of policy iteration

In practice:

- Policy evaluation can be stopped when:

$$\|Q_{\tau+1} - Q_\tau\| \leq \varepsilon_{\text{peval}}$$

- Policy iteration can be stopped when:

$$\|h_{\ell+1} - h_\ell\| \leq \varepsilon_{\text{piter}}$$

- Note:  $\varepsilon_{\text{piter}}$  can be taken 0!

# Q-iteration vs. policy iteration

## Number of iterations to convergence

- Q-iteration > policy iteration

## Complexity

- one iteration of Q-iteration
  - > one iteration of iterative policy evaluation
- complete Q-iteration **???** complete policy iteration

- 1 Introduction
- 2 Markov decision process & optimal solution
- 3 Dynamic programming, DP
- 4 Monte Carlo, MC
- 5 Exploration basics
- 6 Temporal differences, TD
- 7 Improving data efficiency

# Algorithm landscape

By model usage:

- Model-based:  $f, \rho$  known a priori
- Model-free:  $f, \rho$  unknown
- Model-learning:  $f, \rho$  learned from data

By interaction level:

- Offline: algorithm runs in advance
- Online: algorithm runs with the system

We move to online RL for the remainder of the talk



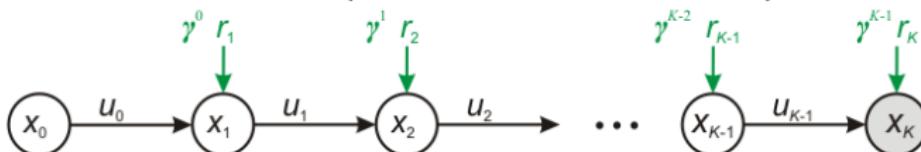
# Policy evaluation change

To find  $Q^h$ :

- So far: model-based policy evaluation
- Reinforcement learning: model not available!
- Learn  $Q^h$  from data obtained by  
**online interaction with the system**

# Monte Carlo policy evaluation

Recall:  $Q^h(x_0, u_0) = \mathbb{E}_{x_1} \{ \rho(x_0, u_0, x_1) + \gamma R^h(x_1) \}$



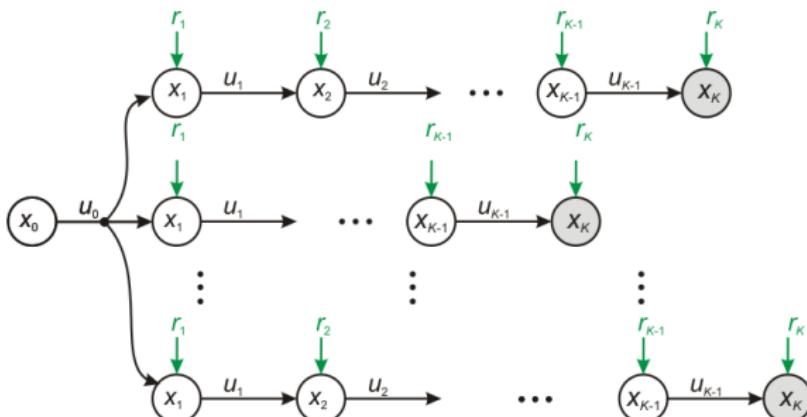
- Trial from  $(x_0, u_0)$  to  $x_K$  using  $u_1 = h(x_1)$ ,  $u_2 = h(x_2)$ , etc.
  - $x_K$  must be terminal (assumed further) or  $K$  large enough
- ⇒ Return along trial provides a sample of  $Q^h(x_0, u_0)$ :

$$\sum_{j=0}^{K-1} \gamma^j r_{j+1}$$

- Furthermore, sample of  $Q^h(x_k, u_k)$  at each step  $k$ :

$$Q^h(x_k, u_k) = \sum_{j=k}^{K-1} \gamma^{j-k} r_{j+1}$$

# Monte Carlo policy evaluation (cont'd)



- To learn the expected value, run  $N$  trajectories (often called roll-outs)
- Estimated Q-value = **average** of the returns, e.g.

$$Q^h(x_0, u_0) = \frac{1}{N} \sum_{i=1}^N \sum_{j=0}^{K_i-1} \gamma^j r_{i,j+1}$$

# Monte Carlo policy iteration

## Monte Carlo policy iteration

```
for each iteration  $\ell$  do
    run  $N$  trials applying  $h_\ell$ 
    reset accumulator  $A(x, u)$ , counter  $C(x, u)$  to 0
    for each step  $k$  of each trial  $i$  do
         $A(x_k, u_k) \leftarrow A(x_k, u_k) + \sum_{j=k}^{K_i-1} \gamma^{j-k} r_{i,j+1}$  (return)
         $C(x_k, u_k) \leftarrow C(x_k, u_k) + 1$ 
    end for
     $Q^{h_\ell}(x, u) \leftarrow A(x, u) / C(x, u)$ 
     $h_{\ell+1}(x) \leftarrow \arg \max_u Q^{h_\ell}(x, u)$ 
end for
```

# Need for exploration

$$Q^h(x, u) \leftarrow A(x, u) / \mathbf{C(x, u)}$$

How to ensure  $C(x, u) > 0$  – **information** about each  $(x, u)$ ?

- ➊ Select representative **initial states**  $x_0$
- ➋ **Actions:**

$u_0$  representative, sometimes different from  $h(x_0)$   
and in addition, perhaps:

$u_k$  representative, sometime different from  $h(x_k)$

# Exploration-exploitation

- **Exploration** needed:  
actions different from the current policy
- **Exploitation** of current knowledge also needed:  
current policy must be applied

Exploration-exploitation dilemma  
– essential in all RL algorithms

(not just in MC)

- 1 Introduction
- 2 Markov decision process & optimal solution
- 3 Dynamic programming, DP
- 4 Monte Carlo, MC
- 5 Exploration basics
- 6 Temporal differences, TD
- 7 Improving data efficiency

# $\varepsilon$ -greedy strategy

- Simple solution to the exploration-exploitation dilemma:  
 **$\varepsilon$ -greedy**

$$u_k = \begin{cases} h(x_k) = \arg \max_u Q(x_k, u) & \text{with probability } (1 - \varepsilon_k) \\ \text{a uniformly random action} & \text{w.p. } \varepsilon_k \end{cases}$$

- Exploration probability  $\varepsilon_k \in (0, 1)$   
usually decreased over time
- Main disadvantage: when exploring, actions are fully random, leading to poor performance

# Softmax strategy

- Action selection:

$$u_k = u \text{ w.p. } \frac{e^{Q(x_k, u)/\tau_k}}{\sum_{u'} e^{Q(x_k, u')/\tau_k}}$$

where  $\tau_k > 0$  is the **exploration temperature**

- Taking  $\tau \rightarrow 0$ , greedy selection recovered;  
 $\tau \rightarrow \infty$  gives uniform random
- Compared to  $\varepsilon$ -greedy, better actions are more likely to be applied even when exploring

# Bandit-based exploration



At single state, exploration modeled as **multi-armed bandit**:

- Action  $j$  = arm with reward distribution  $\rho_j$ , expectation  $\mu_j$
- Best arm (optimal action) has expected value  $\mu^*$
- At step  $k$ , we pull arm (try action)  $j_k$ , getting  $r_k \sim \rho_{j_k}$
- **Objective:** After  $n$  pulls, small regret:  $\sum_{k=1}^n \mu^* - \mu_{j_k}$

# UCB algorithm

Popular algorithm: after  $n$  steps, pick arm with largest **upper confidence bound**:

$$b(j) = \hat{\mu}_j + \sqrt{\frac{3 \log n}{2n_j}}$$

where:

- $\hat{\mu}_j$  = mean of rewards observed for arm  $j$  so far
- $n_j$  = how many times arm  $j$  was pulled

These are only a few simple methods, many others exist,  
e.g. Bayesian exploration, intrinsic rewards etc.

- 1 Introduction
- 2 Markov decision process & optimal solution
- 3 Dynamic programming, DP
- 4 Monte Carlo, MC
- 5 Exploration basics
- 6 Temporal differences, TD
- 7 Improving data efficiency

# DP perspective

- ① Start from policy evaluation:

$$Q_{\tau+1}(x, u) \leftarrow \sum_{x'} f(x, u, x') [\rho(x, u, x') + \gamma Q_\tau(x', h(x'))]$$

- ② Instead of model, use **transition sample** at each step  $k$ ,

$(x_k, u_k, x_{k+1}, r_{k+1}, u_{k+1})$ :

$$Q(x_k, u_k) \leftarrow r_{k+1} + \gamma Q(x_{k+1}, u_{k+1})$$

Note:

$$x_{k+1} \sim f(x_k, u_k, \cdot), r_{k+1} = \rho(x_k, u_k, x_{k+1}), u_{k+1} \sim h(x_{k+1})$$

- ③ Turn into **incremental update**:

$$Q(x_k, u_k) \leftarrow Q(x_k, u_k) + \alpha_k \cdot$$

$$[r_{k+1} + \gamma Q(x_{k+1}, u_{k+1}) - Q(x_k, u_k)]$$

$\alpha_k \in (0, 1]$  learning rate



# Intermediate algorithm

## Temporal differences for policy evaluation

**for** each trial **do**

  init  $x_0$ , choose initial action  $u_0$

**repeat** at each step  $k$

    apply  $u_k$ , measure  $x_{k+1}$ , receive  $r_{k+1}$

    choose **next** action  $u_{k+1} \sim h(x_{k+1})$

$$Q(x_k, u_k) \leftarrow Q(x_k, u_k) + \alpha_k \cdot$$

$$[r_{k+1} + \gamma Q(x_{k+1}, u_{k+1}) - Q(x_k, u_k)]$$

**until** trial finished

**end for**

# MC perspective

Temporal differences for policy  $h$  evaluation

**for** each trial **do**

...

**repeat** each step  $k$

    apply  $u_k$ , measure  $x_{k+1}$ , receive  $r_{k+1}$

$Q(x_k, u_k) \leftarrow \dots Q\dots$

**until** trial finished

**end for**

Monte Carlo

**for** each trial **do**

    execute trial

...

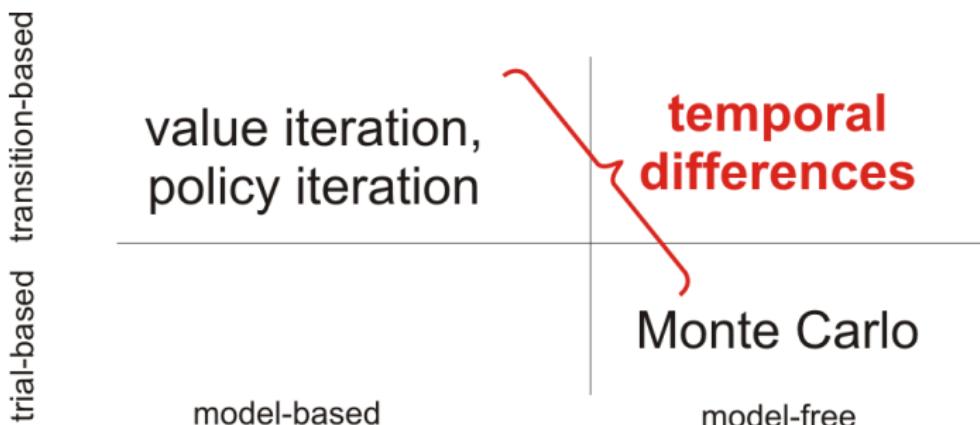
$Q(x, u) \leftarrow A(x, u)/C(x, u)$

**end for**



# MC and DP perspectives

- Learn from online interaction: like MC, unlike DP
- Update after each transition, using previous Q-values:  
like DP, unlike MC



# Exploration-exploitation

choose next action  $u_{k+1} \sim h(x_{k+1})$

- Information about  $(x, u) \neq (x, h(x))$  needed  
⇒ **exploration**
- $h$  must be followed  
⇒ **exploitation**

# Policy improvement

- Previous algorithm:  $h$  fixed
- Improving  $h$ : simplest, after each transition,  
called optimistic policy improvement
- $h$  implicit, greedy in  $Q$   
(update  $Q \Rightarrow$  implicitly improve  $h$ )
- E.g.  $\varepsilon$ -greedy:

$$u_{k+1} = \begin{cases} \arg \max_u Q(x_{k+1}, u) & \text{w.p. } (1 - \varepsilon_{k+1}) \\ \text{uniformly random} & \text{w.p. } \varepsilon_{k+1} \end{cases}$$

# SARSA

## SARSA

**for** each trial **do**

    init  $x_0$

    choose  $u_0$  with exploration based on  $Q$

**repeat** at each step  $k$

        apply  $u_k$ , measure  $x_{k+1}$ , receive  $r_{k+1}$

        choose  $u_{k+1}$  with exploration based on  $Q$

$$Q(x_k, u_k) \leftarrow Q(x_k, u_k) + \alpha_k \cdot$$

$$[r_{k+1} + \gamma Q(x_{k+1}, u_{k+1}) - Q(x_k, u_k)]$$

**until** trial finished

**end for**

Origin of the name:  $(x_k, u_k, r_{k+1}, x_{k+1}, u_{k+1}) =$   
**(State, Action, Reward, State, Action)**

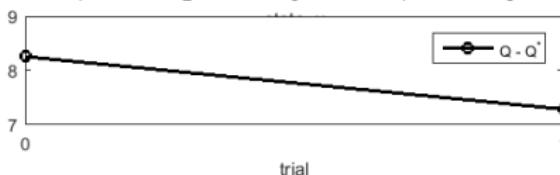
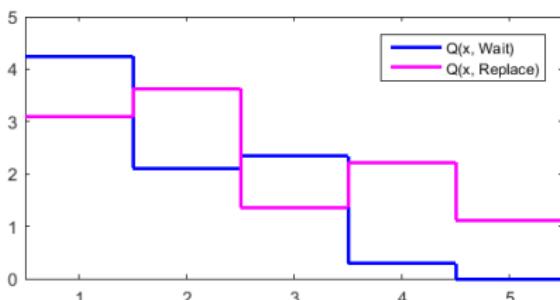


# Machine replacement: SARSA demo

Parameters:  $\alpha = 0.1$ ,  $\varepsilon = 0.3$  (constant), single trial

$x_0 = 1$

SARSA, trial 1 completed



# Q-learning

- 1 Similarly to SARSA, start from Q-iteration:

$$Q_{\ell+1}(x, u) \leftarrow \sum_{x'} f(x, u, x')[\rho(x, u, x') + \gamma \max_{u'} Q_\ell(x', u')]$$

- 2 Instead of model, use at each step  $k$  **transition sample**

$(x_k, u_k, x_{k+1}, r_{k+1})$ :

$$Q(x_k, u_k) \leftarrow r_{k+1} + \gamma \max_{u'} Q(x_{k+1}, u')$$

Note:  $x_{k+1} \sim f(x_k, u_k, \cdot)$ ,  $r_{k+1} = \rho(x_k, u_k, x_{k+1})$

- 3 Turn into **incremental** update:

$$Q(x_k, u_k) \leftarrow Q(x_k, u_k) + \alpha_k \cdot$$

$$[r_{k+1} + \gamma \max_{u'} Q(x_{k+1}, u') - Q(x_k, u_k)]$$



# Q-learning

## Q-learning

**for** each trial **do**

  init  $x_0$

**repeat** at each step  $k$

    choose  $u_k$  with exploration based on  $Q$

    apply  $u_k$ , measure  $x_{k+1}$ , receive  $r_{k+1}$

$$Q(x_k, u_k) \leftarrow Q(x_k, u_k) + \alpha_k \cdot$$

$$[r_{k+1} + \gamma \max_{u'} Q(x_{k+1}, u') - Q(x_k, u_k)]$$

**until** trial finished

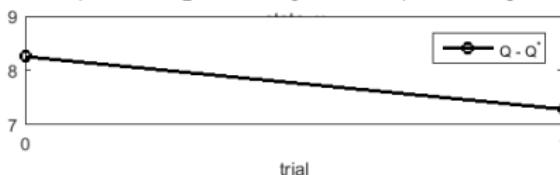
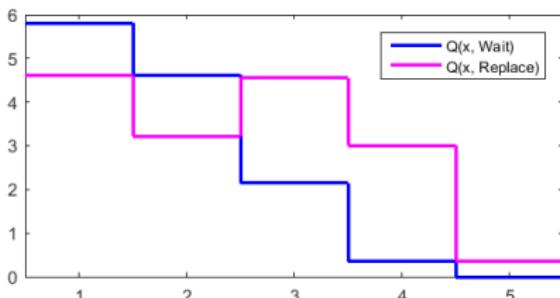
**end for**



# Machine replacement: Q-learning demo

Parameters:  $\alpha = 0.1$ ,  $\varepsilon = 0.3$  (constant), single trial

$x_0 = 1$



# Convergence

Conditions for convergence to  $Q^*$   
in both SARSA and Q-learning:

- ➊ All pairs  $(x, u)$  continue to be updated:  
requires **exploration**, e.g.  $\varepsilon$ -greedy
- ➋ Technical conditions on  $\alpha_k$  (goes to 0,  $\sum_{k=0}^{\infty} \alpha_k^2 = \text{finite}$ ,  
but not too fast,  $\sum_{k=0}^{\infty} \alpha_k \rightarrow \infty$ )

In addition, for SARSA:

- ➌ Policy must become greedy asymptotically  
e.g. for  $\varepsilon$ -greedy,  $\lim_{k \rightarrow \infty} \varepsilon_k = 0$

# Discussion

## SARSA **on-policy**

- Always updates towards the Q-function of the current policy

## Q-learning **off-policy**

- Irrespective of the current policy, always updates towards optimal Q-function

## Advantages of temporal differences

- Easy to understand and implement
- Low complexity  $\Rightarrow$  fast execution

Exploration and  $\alpha_k$  sequence **greatly influence** performance

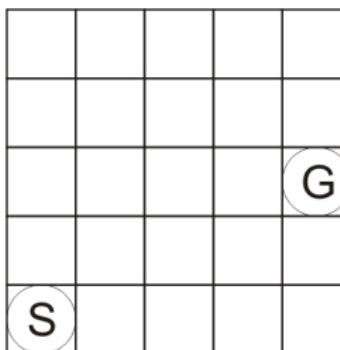


- 1 Introduction
- 2 Markov decision process & optimal solution
- 3 Dynamic programming, DP
- 4 Monte Carlo, MC
- 5 Exploration basics
- 6 Temporal differences, TD
- 7 Improving data efficiency

# Motivation

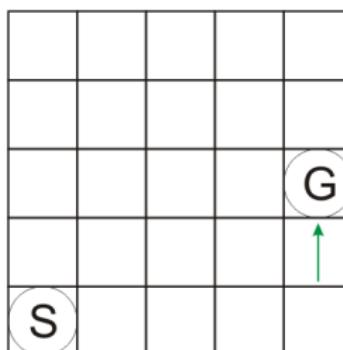
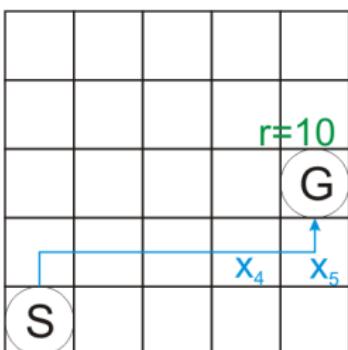
TD uses data inefficiently, and data is often costly.

Example:



- 2D gridworld navigation from **Start** to **Goal**
- Nonzero reward = **10** only on reaching G (terminal state)

## Motivation (cont'd)



- Take SARSA with  $\alpha = 1$ ; initialize  $Q = 0$
- Updates along the trial on the left:

...

$$Q(x_4, u_4) = 0 + \gamma \cdot Q(x_5, u_5) = 0$$

$$Q(x_5, u_5) = 10 + \gamma \cdot 0 = 10$$

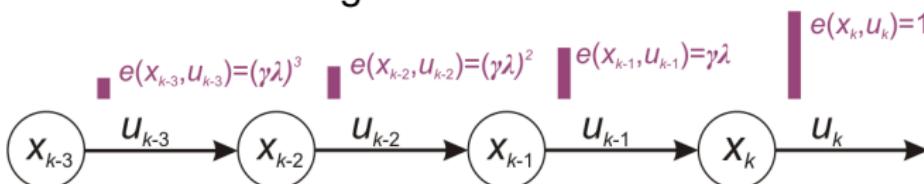
- A new transition from  $x_4$  to  $x_5$  (and hence a new trial) required to propagate the info to  $x_4$ !

# Ideas presented

- ➊ Eligibility traces
- ➋ Experience replay

# Eligibility traces

- Idea: Leave a **trace** along the trial:



- $\lambda \in [0, 1]$  decay rate

## Replacing traces

```
 $e(x, u) \leftarrow 0$  for all  $x, u$ 
for each step  $k$  do
     $e(x, u) \leftarrow \lambda\gamma e(x, u)$  for all  $x, u$ 
     $e(x_k, u_k) \leftarrow 1$ 
end for
```

## Example algorithm: SARSA( $\lambda$ )

- Recall original SARSA only updates  $Q(x_k, u_k)$ :

$$Q(x_k, u_k) \leftarrow Q(x_k, u_k) + \alpha_k \cdot$$

$$[r_{k+1} + \gamma Q(x_{k+1}, u_{k+1}) - Q(x_k, u_k)]$$

- SARSA( $\lambda$ ) updates **all eligible pairs**:

$$Q(x, u) \leftarrow Q(x, u) + \alpha_k \cdot e(x, u) \cdot$$

$$[r_{k+1} + \gamma Q(x_{k+1}, u_{k+1}) - Q(x_k, u_k)] \quad \forall x, u$$

# SARSA( $\lambda$ )

## SARSA( $\lambda$ )

**for** each trial **do**

  init  $x_0$

$e(x, u) \leftarrow 0 \quad \forall x, u$

  choose  $u_0$  with exploration based on  $Q$

**repeat** at each step  $k$

    apply  $u_k$ , measure  $x_{k+1}$ , receive  $r_{k+1}$

    choose  $u_{k+1}$  with exploration based on  $Q$

$e(x, u) \leftarrow \lambda \gamma e(x, u) \quad \forall x, u$

$e(x_k, u_k) \leftarrow 1$

$Q(x, u) \leftarrow Q(x, u) + \alpha_k \cdot e(x, u) \cdot$

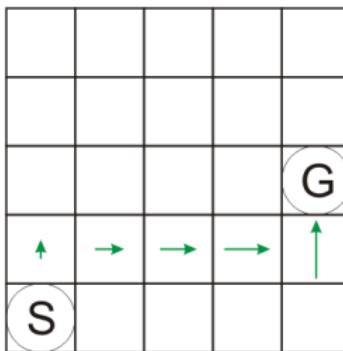
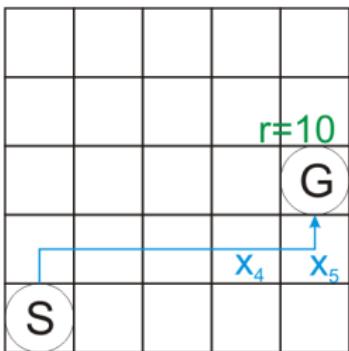
$[r_{k+1} + \gamma Q(x_{k+1}, u_{k+1}) - Q(x_k, u_k)]$  for all  $x, u$

**until** trial finished

**end for**



## Example: Effect of eligibility traces



- $\lambda = 0.7$
- Updates until  $x_4$ :  $Q$  remains 0
- At  $x_5$ , the entire trial gets updated:

$$Q(x_5, u_5) = 10 + \gamma 0 = 10$$

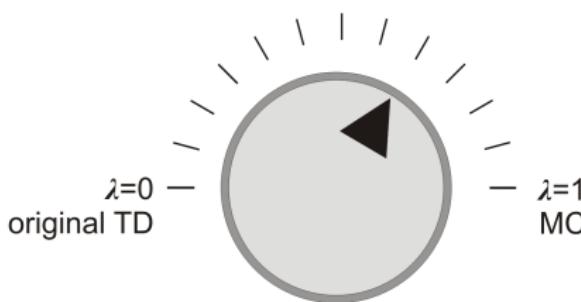
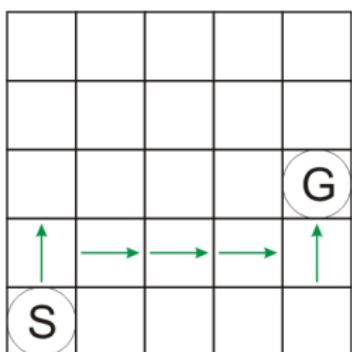
$$Q(x_4, u_4) = (\gamma\lambda)[10 + \gamma 0] = 3.5$$

$$Q(x_3, u_3) = (\gamma\lambda)^2[10 + \gamma 0] = 1.225$$

...

# TD versus MC

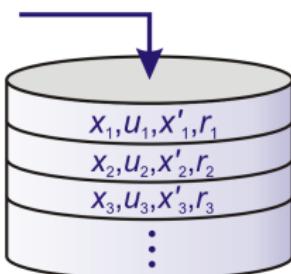
- $\lambda = 0 \Rightarrow$  recovers original algorithms, e.g. SARSA(0)
- $\lambda = 1 \Rightarrow$  TD becomes like MC



Typical values of  $\lambda$  are around 0.5 to 0.8

# Experience replay

- Store each transition  $(x_k, u_k, x'_{k+1}, r_{k+1})$  in a database



- At each step, **replay**  $N$  transitions from the database (in addition to regular updates)

# Q-learning with experience replay

## Q-learning with experience replay

**for** each trial **do**

    init  $x_0$

**repeat** at each step  $k$

        apply  $u_k$ , measure  $x_{k+1}$ , receive  $r_{k+1}$

$$Q(x_k, u_k) \leftarrow Q(x_k, u_k) + \alpha_k \cdot$$

$$[r_{k+1} + \gamma \max_{u'} Q(x_{k+1}, u') - Q(x_k, u_k)]$$

        add  $(x_k, u_k, x_{k+1}, r_{k+1})$  to database

        ReplayExperience

**until** trial finished

**end for**

# ReplayExperience procedure

## ReplayExperience

**loop**  $N$  times

retrieve a transition  $(x, u, x', r)$  from database

$$Q(x, u) \leftarrow Q(x, u) + \alpha \cdot$$

$$[r + \gamma \max_{u'} Q(x', u') - Q(x, u)]$$

**end loop**

Retrieval order:

- ① Backwards along trials, best for classical algos
- ② Randomly, helps e.g. in deep RL
- ③ etc.



# Textbooks

- Sutton & Barto, *Reinforcement Learning: An Introduction*, 1998 (+ ongoing 2nd edition, 2017).
- Bertsekas, *Dynamic Programming and Optimal Control*, vol. 2, 4th ed., 2012.
- Szepesvári, *Algorithms for Reinforcement Learning*, 2010.
- Buşoniu, Babuška, De Schutter, & Ernst, *Reinforcement Learning and Dynamic Programming Using Function Approximators*, 2010.

